



Tutorial on QuantumFlow+VACSEN: A Visualization System for Quantum Neural Networks on Noisy Quantum Devices

Session 4: Quantum Neural Network Compression

Zhepeng Wang

Ph.D. Student

Electrical and Computer Engineering

George Mason University

zwang48@gmu.edu

<https://jqub.ece.gmu.edu>

How to Compress a Quantum Neural Network?

Quantum Neural Network Compression

<https://arxiv.org/pdf/2207.01578.pdf>


Zhepeng Wang (Presenter), Zhirui Hu, Dr. Weiwen Jiang

Department of Electrical and Computer Engineering

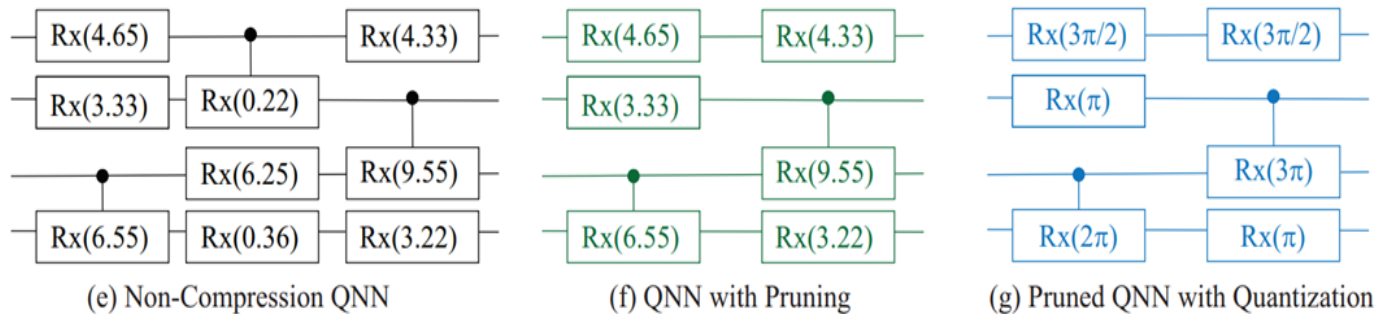
Jqub @ George Mason University

Motivation and Background

- Pruning and Quantization in Classical ML



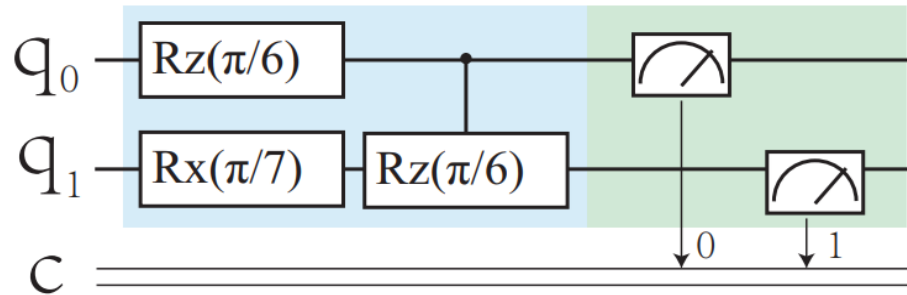
- Pruning and Quantization in Quantum ML



- Pruning:** Not only 0 can be pruned, but also 2π , 4π , etc.
- Quantization:** Different quantization level may have different cost

Motivation and Background

- Quantum Neural Network Compression Should be Compilation Aware



Compilation

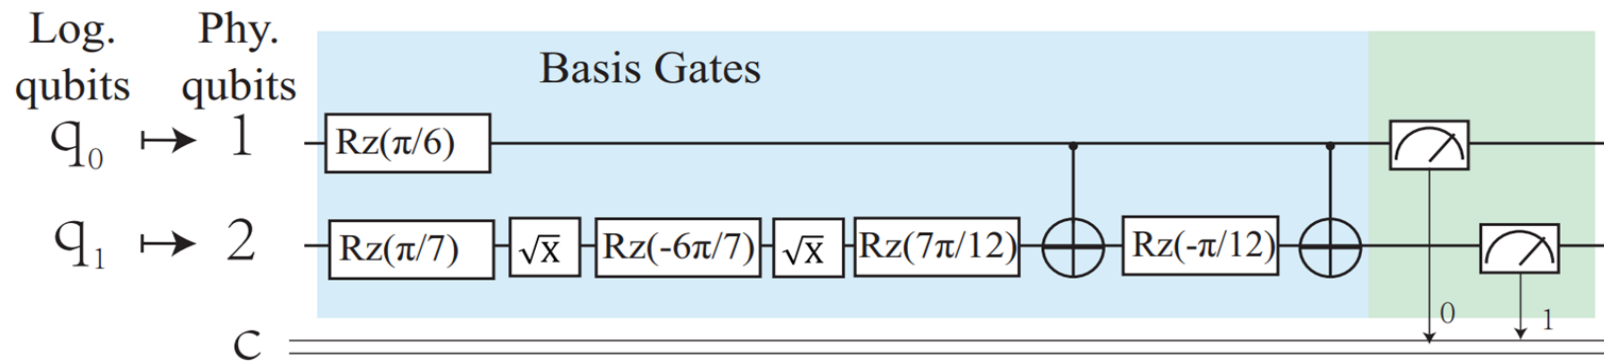


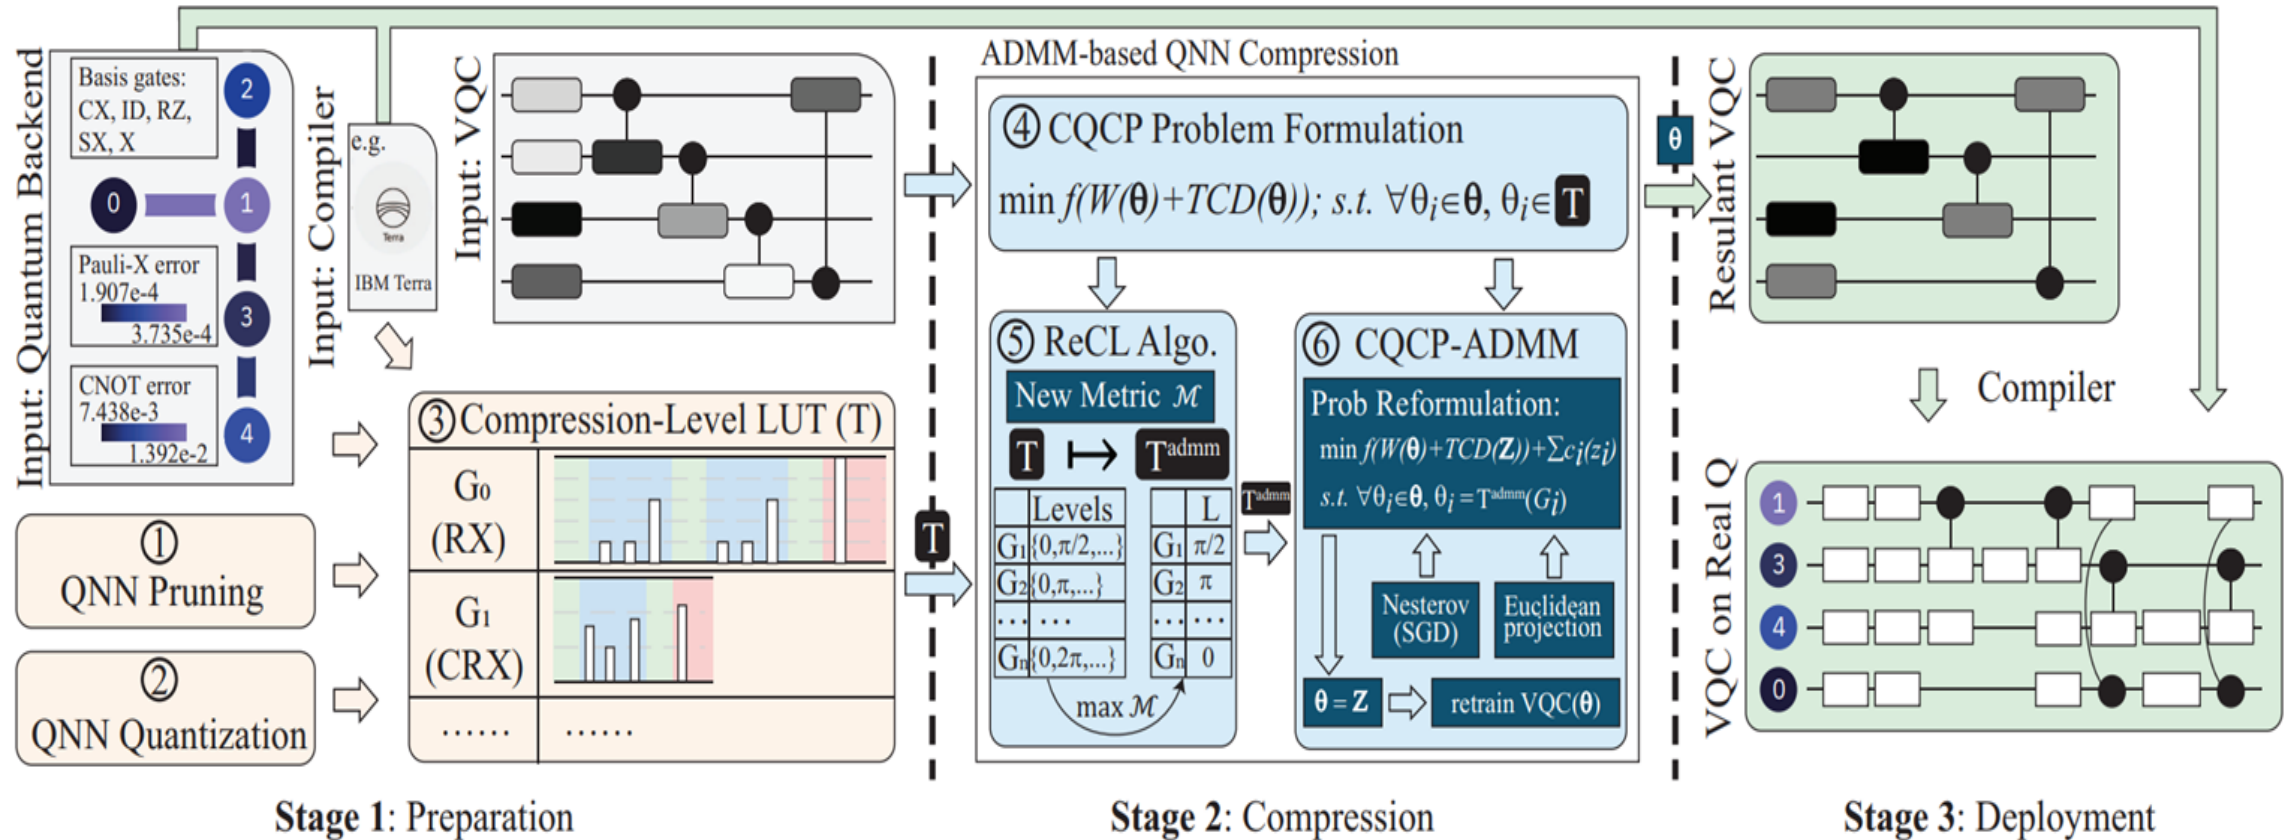
Table 1: circuit depth of compiled quantum gates on IBM quantum processors; parameters are in the range of $[0, 4\pi]$

Gate	0	π	2π	3π	4π	$\pi/2$	$3\pi/2$	$5\pi/2$	$7\pi/2$	others
RX	0	1	0	1	0	1	3	1	3	5
RY	0	2	0	2	0	3	3	3	3	4
CRX	0	8	5	9	0	11	11	11	11	11
CRY	0	8	6	8	0	10	10	10	10	10

CompVQC

- General Overview

Three stages: 1. Preparation; 2. Compression; 3. Deployment

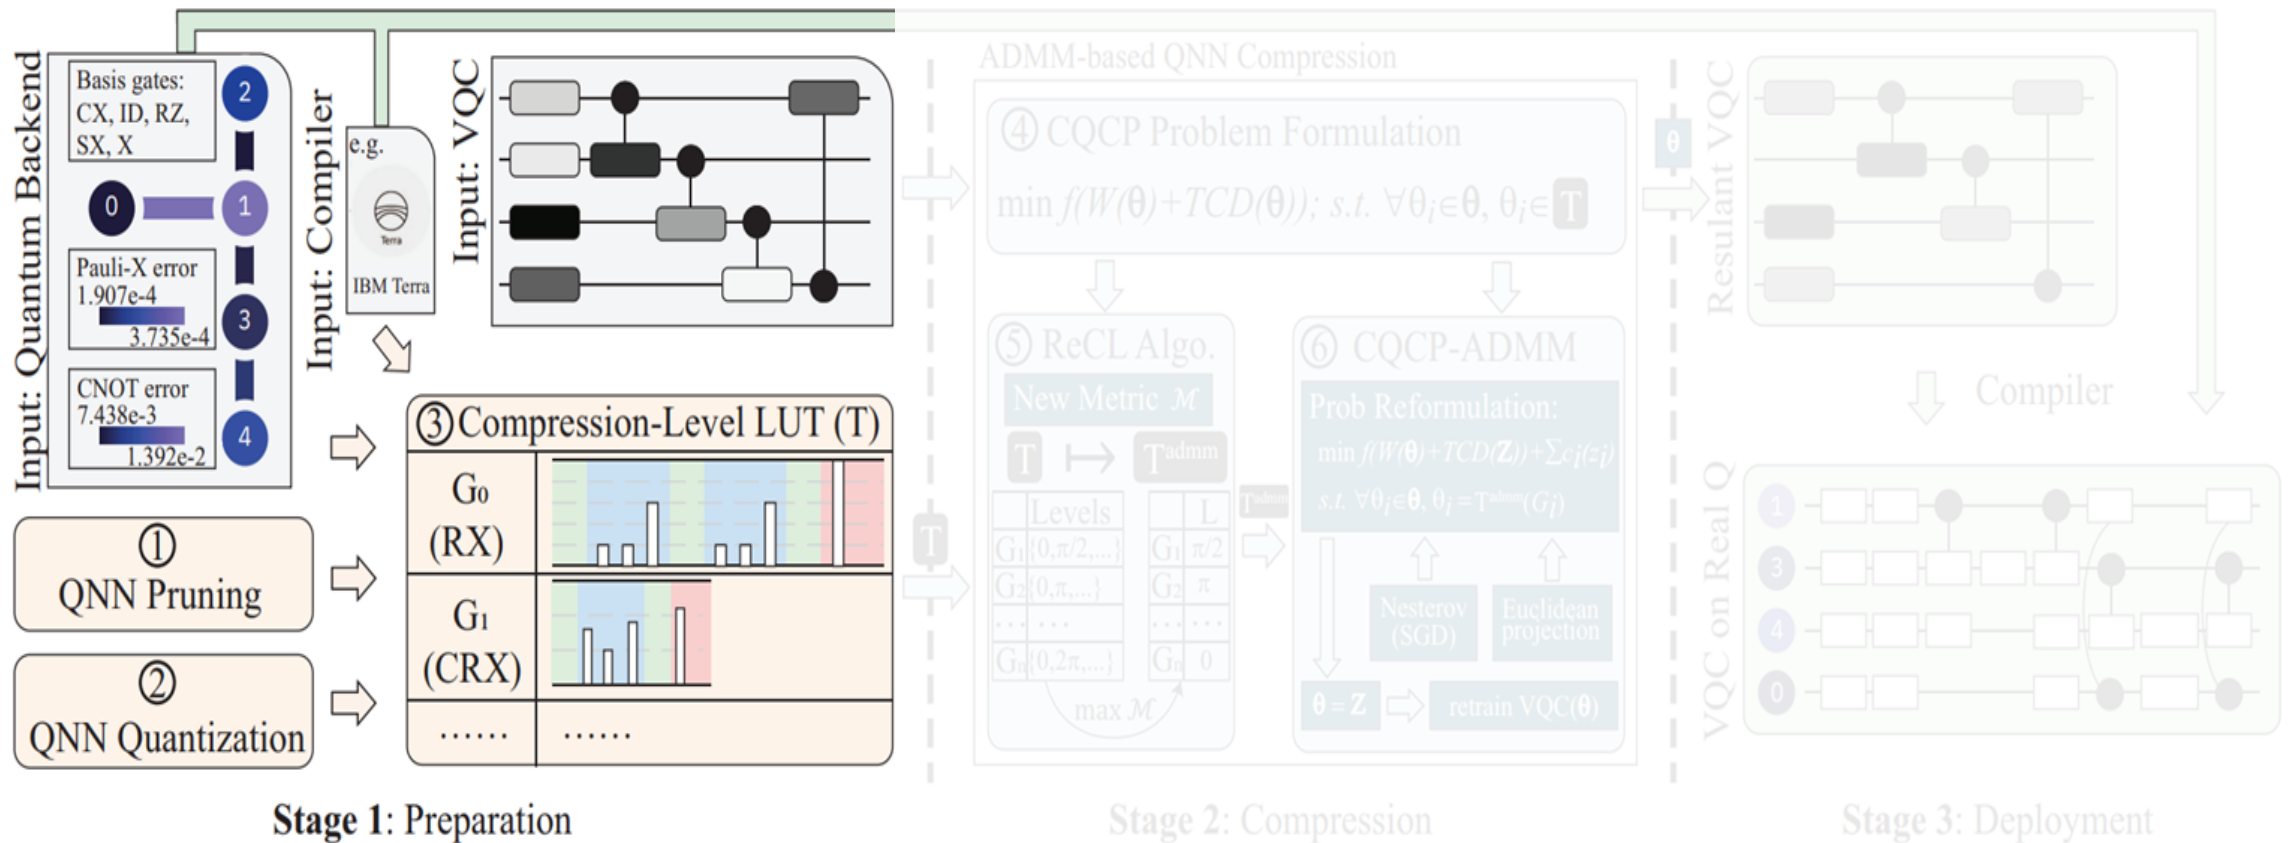


CompVQC

- LUT Construction and Training a Quantum Model
- Reconstruct LUT for ADMM
- Compression based on ADMM
- Deployment

CompVQC

- LUT Construction and Training a Quantum Model



CompVQC

- LUT Construction and Training a Quantum Model

- ❑ **Compression-Level Lookup Table (LUT)**

A combination of pruning/quantization level called as “compression level”.

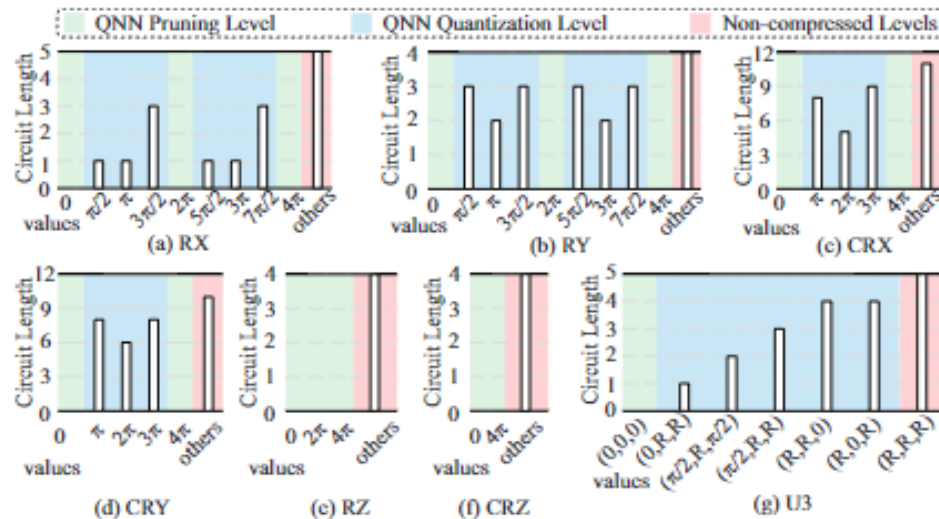


Table 1: circuit depth of compiled quantum gates on IBM quantum processors; parameters are in the range of $[0, 4\pi]$

Gate	0	π	2π	3π	4π	$\pi/2$	$3\pi/2$	$5\pi/2$	$7\pi/2$	others
RX	0	1	0	1	0	1	3	1	3	5
RY	0	2	0	2	0	3	3	3	3	4
CRX	0	8	5	9	0	11	11	11	11	11
CRY	0	8	6	8	0	10	10	10	10	10

- ❑ **VQC Pre-Training**

A VQC model is pre-trained for compression and the training process is implemented with **Torch Quantum**.

Hands-On Tutorial (1) : LUT Construction

Input

- Fixing points list
- Logical Gates List to be used
- Quantum Backend

Do

- Get the compiler for the backend
- Get the compiled circuit length of each logical gate at each special fixing points

Output

- Get the compiler for the backend

```
#Input
test_fixing_points = [math.pi*4, math.pi*2, math.pi, math.pi*3, math.pi/2,
                     math.pi/2*5, math.pi/2*7, math.pi/2*3, math.pi/6]
logical_gates = ['rx', 'ry', 'rz', 'crx', 'cry', 'crz']
backend = FakeValencia()

#api
df = LUT_construction(test_fixing_points, logical_gates, backend)
```

fixing_points	rx	ry	rz	crx	cry	crz
12.57	0	0	0	0	0	0
6.28	0	0	0	5	6	4
3.14	1	2	1	8	8	4
9.42	1	2	1	9	8	4
1.57	1	3	1	11	10	4
7.85	1	3	1	11	10	4
11.00	3	3	1	11	10	4
4.71	3	3	1	11	10	4
0.52	5	4	1	11	10	4

Hands-On Tutorial (1)

LUT Construction



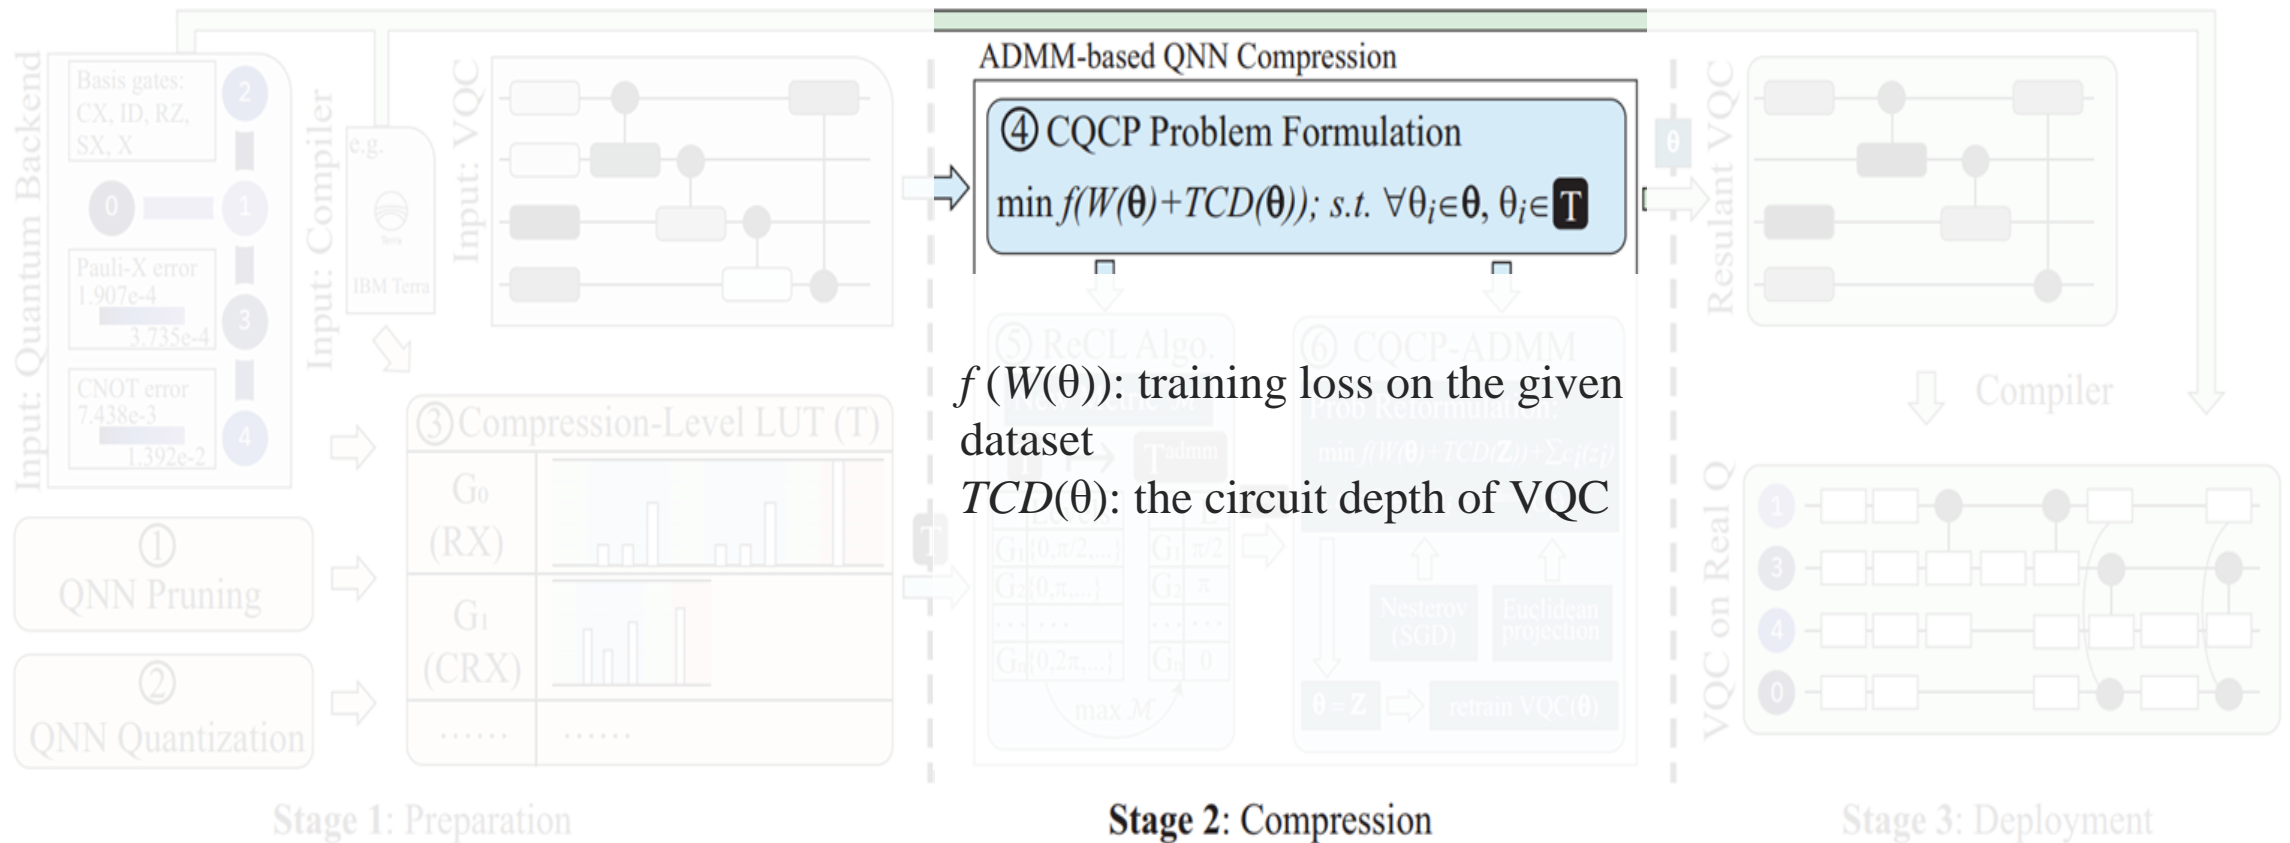
CompVQC

- LUT Construction and Training a Quantum Model
- **Reconstruct LUT for ADMM**
- Compression based on ADMM
- Deployment

CompVQC

- Problem Definition

Given VQC $W(\theta)$, LUT T , quantum compiler C , the problem is to determine trainable parameters θ , such that:



CompVQC

- Reconstruction LUT for ADMM

Process is conducted by traversing all quantum gates in VQC and select the compression target with highest metric.

A heuristic metric for the choice

$$\mathcal{M}(\theta, G_i(\gamma_{i,k})) = \text{acc}(W(\theta^{i,k})) \cdot \tau(\theta^{i,k}, \theta)$$

$$\tau(\theta^{i,k}, \theta) = \frac{TCD(\theta)}{TCD(\theta^{i,k})}$$

$\text{acc}(W(\theta^{i,k}))$: the accuracy of the VQC under new parameters

$TCD(\theta)$: the inverse of the compression ratio by changing parameters from θ to $\theta^{i,k}$

ADMM-based QNN Compression

④ CQCP Problem Formulation

$$\min f(W(\theta) + TCD(\theta)); \text{ s.t. } \forall \theta_i \in \theta, \theta_i \in \mathcal{T}$$

⑤ ReCL Algo.

New Metric \mathcal{M}

$$\mathcal{T} \mapsto \mathcal{T}^{\text{admm}}$$

Levels	L
$G_1\{0, \pi/2, \dots\}$	$G_1 \pi/2$
$G_2\{0, \pi, \dots\}$	$G_2 \pi$
\dots	\dots
$G_n\{0, 2\pi, \dots\}$	$G_n 0$

max \mathcal{M}

⑥ CQCP-ADMM

Prob Reformulation:

$$\min f(W(\theta) + TCD(\mathbf{Z})) + \sum c_j(z_j) \\ \text{ s.t. } \forall \theta_i \in \theta, \theta_i = \mathcal{T}^{\text{admm}}(G_i)$$

Nesterov (SGD)

Euclidean projection

$$\theta = \mathbf{Z}$$

retrain VQC(θ)

Resulant VQC

VQC on Real Q

Compiler

Stage 1: Preparation

Stage 2: Compression

Stage 3: Deployment

Hands-On Tutorial (2) : Reconstruct LUT for ADMM

Input

- trained model
- Original LUT
- The metrics function of accuracy and length

```
#input
model = torch.load('model.pth')
lut = pd.read_csv('lut.csv')
def metrics_func(acc, depth):
    return acc+1.0/depth
backend = FakeValencia()
```

For each parameter, Do

- Replace it with points at compression level in original LUT while fixing other parameters
- Calculate the metrics of each new model
- Select the point with the highest metric as the compression level for ADMM

```
#api
new_lut = LUT_reconstruction(model, lut, backend, metrics_func)
```

Output

- A new LUT for ADMM

```
[ 7.85  1.57 11.   0.   3.14  1.57  3.14  3.14  0.   0.   0.   0.
  0.   0.   0.   0.   0.   0. ]
```

Hands-On Tutorial (2)

Reconstruct LUT for ADMM



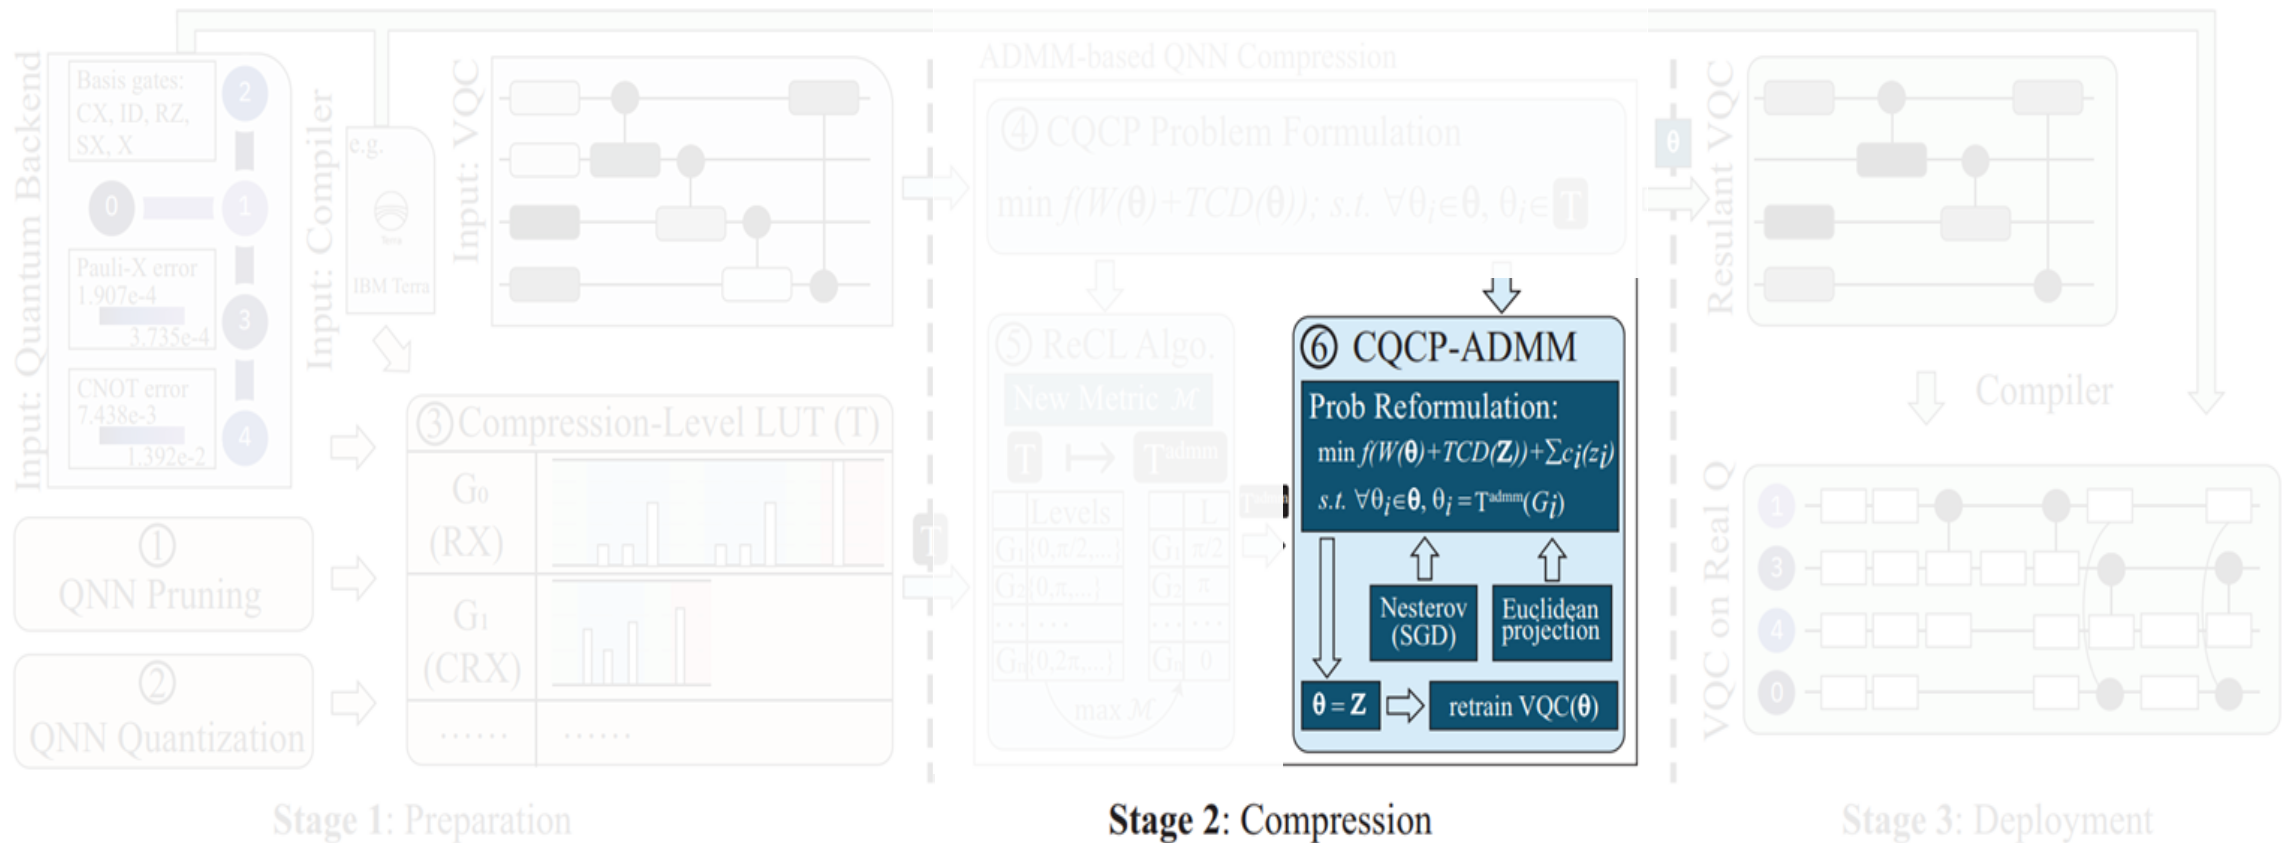
CompVQC

- LUT Construction and Training a Quantum Model
- Reconstruct LUT for ADMM
- **Compression based on ADMM**
- Deployment

CompVQC

- Compression based on ADMM

Each parameter can either be compressed to the target value in T^{admm} or not compressed.



CompVQC

- Compression based on ADMM

Given reconstructed compression-level LUT T^{admm} , the CQCP is formulated as:

$$\begin{aligned} \min_{\{\theta_i\}} \quad & f(W(\theta)) + TCD(Z) + \sum_{\forall z_i \in Z} c_i(Z_i), \\ \text{s.t.} \quad & \forall \theta_i \in \theta, \quad \theta_i = T^{admm}(G_i). \end{aligned}$$

Z : a set of auxiliary variables for subproblem decomposition and $z_i \in Z$ is corresponding to $\theta_i \in \theta$
 $f(W(\theta)) + TCD(Z)$: the objective function in the original CQCP problem(Previously seen).

$$c_i(Z_i) = \begin{cases} 0 & \text{if } \theta_i \in T^{s,r}(G_i), T^{s,r} = T^{admm} \odot \text{mask}^r \\ +\infty & \text{if otherwise.} \end{cases}$$

$c_i(Z_i)$: An indicator function to serve as a penalty term

mask^r : variable to indicate whether the parameters will be compressed at iteration r .

Hands-On Tutorial (3) : Compression based on ADMM

Input

- A trained model
- A new LUT for ADMM

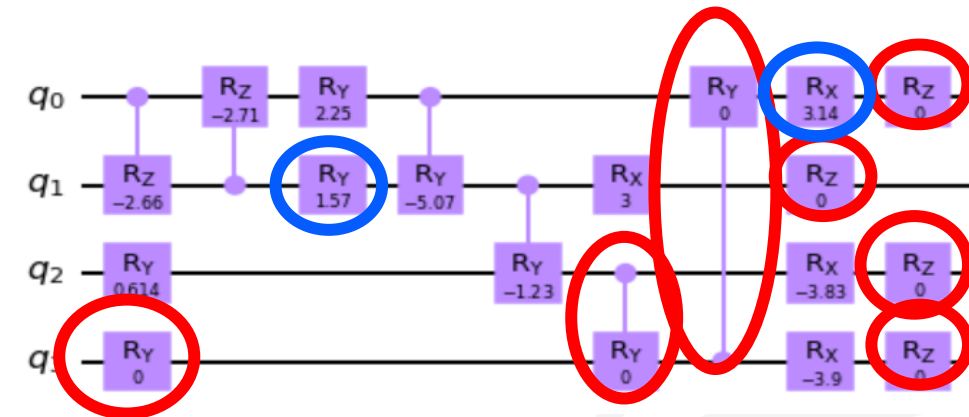

Do

- Compress a model with ADMM
- Fine-tune the compressed model

Output

- A compressed model

	Circuit Length	Accuracy
Original model	51	94.2%
Compressed	35	97.10%



Pruned

Quantized

Hands-On Tutorial (3)

Compression based on ADMM

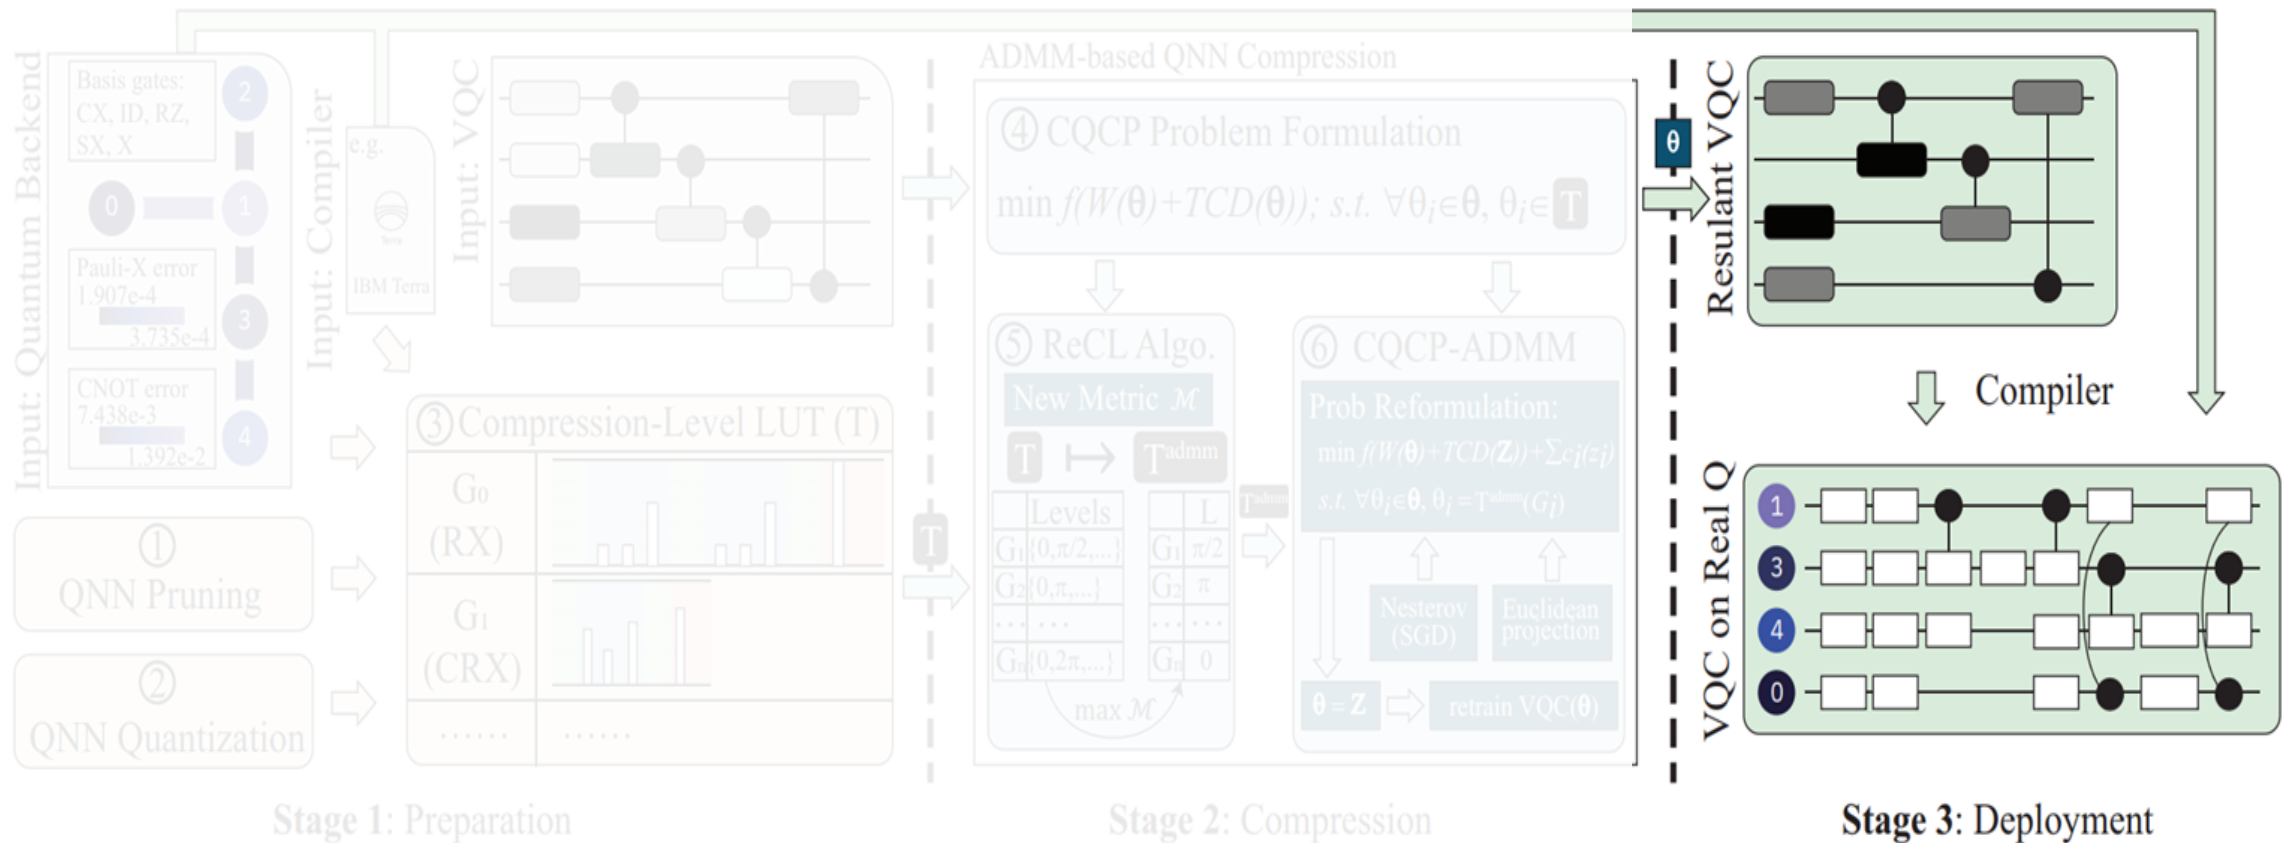


CompVQC

- LUT Construction and Training a Quantum Model
- Reconstruct LUT for ADMM
- Compression based on ADMM
- **Deployment**

CompVQC

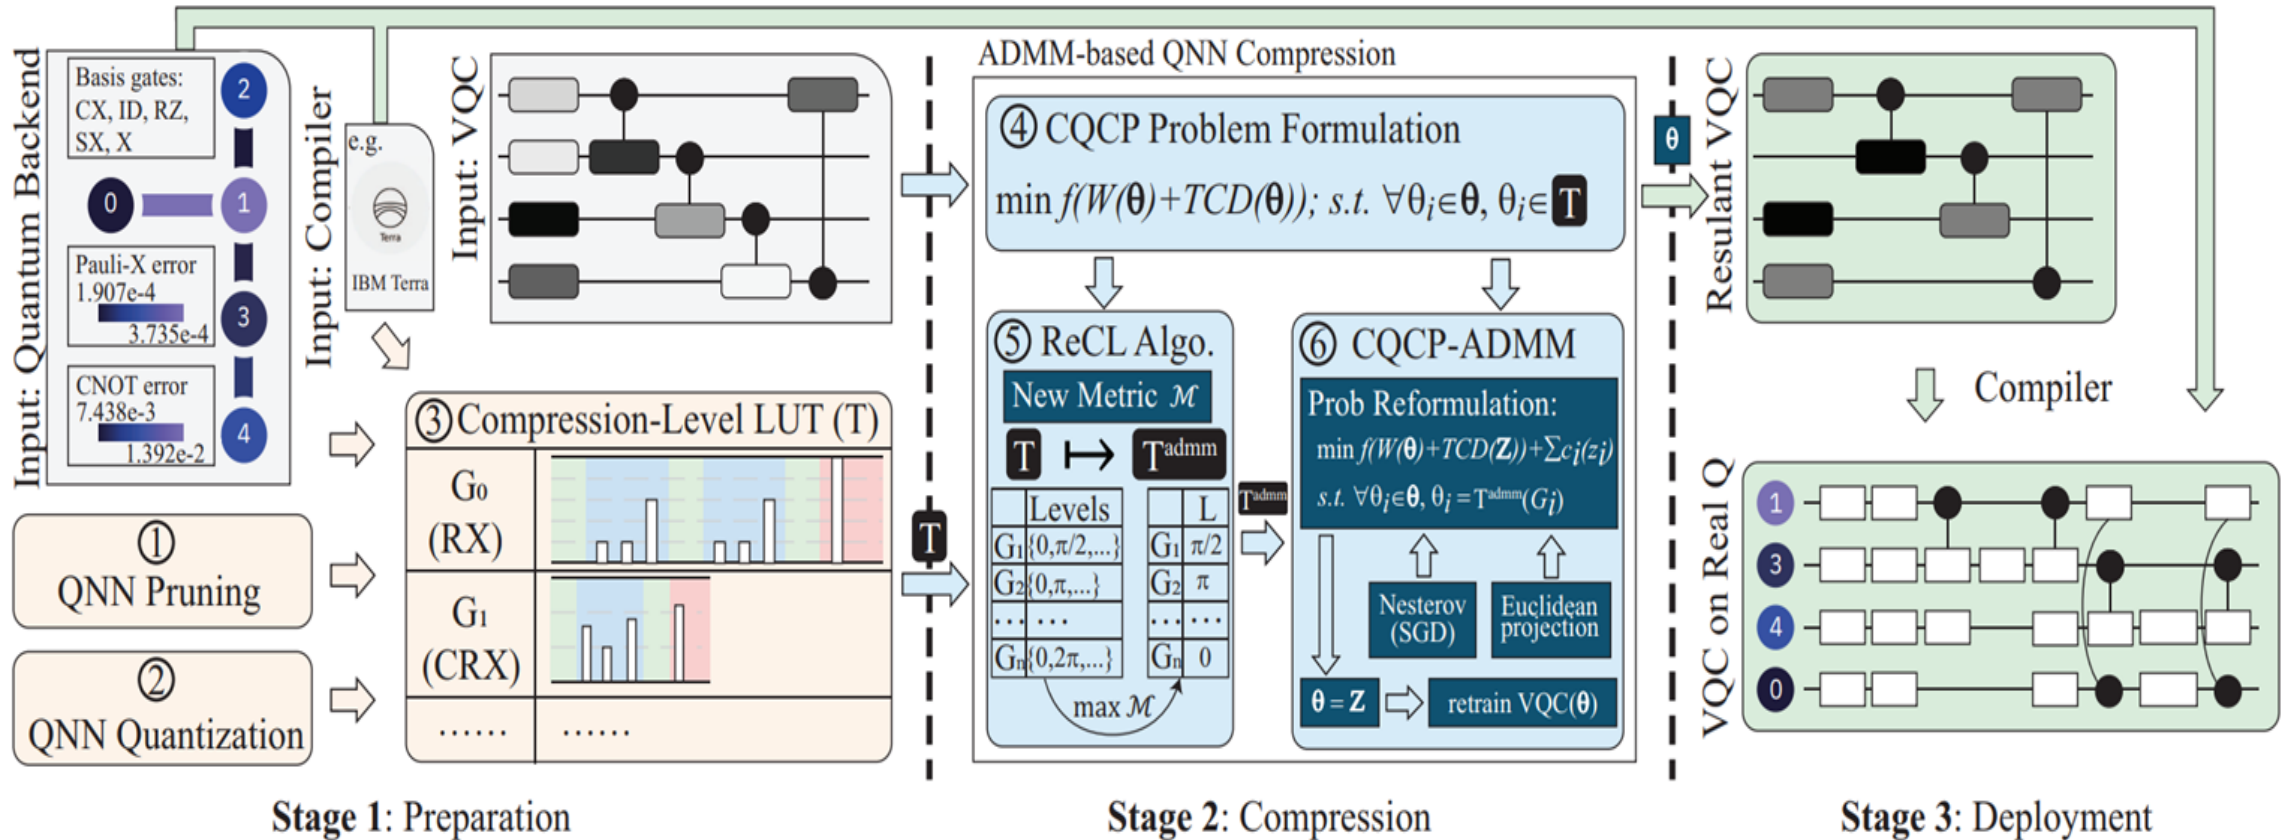
- Deployment



CompVQC

- General Overview

Three stages: 1. Preparation; 2. Compression; 3. Deployment



Experimental Results

- Simulation Results on ML Dataset

CompVQC can maintain high accuracy with **<1% accuracy loss**. And the reduction of circuit length is up to **2.5X**.

Table 2: Comparison among different methods on the accuracy performance and the TCD of the VQC

Compression Method	MNIST-2		Fashion-MNIST-2	
	Acc. (vs. Baseline)	TCD (Speedup)	Acc. (vs. Baseline)	TCD (Speedup)
Vanilla VQC	82.74%(0)	121(0)	87.58%(0)	92(0)
Zero-Only-Pruning	80.58%(-2.16%)	70(1.73×)	86.92%(-0.67%)	63(1.46×)
CompVQC-Pruning	81.83%(-0.91%)	74(1.64 ×)	87.41%(-0.17%)	47(1.96×)
CompVQC-Quant	80.99%(-1.75%)	108(1.10×)	86.25%(-1.33%)	74(1.24×)
CompVQC	81.83%(-0.91%)	47(2.57×)	87.58%(-0.00%)	47(1.96×)




Figure 5: Main results: The Accuracy-Circuit Depth Tradeoff on Fashion-MNIST2




Figure 6: Main Results: CompVQC Scalability on Fashion-MNIST with 2-4 class

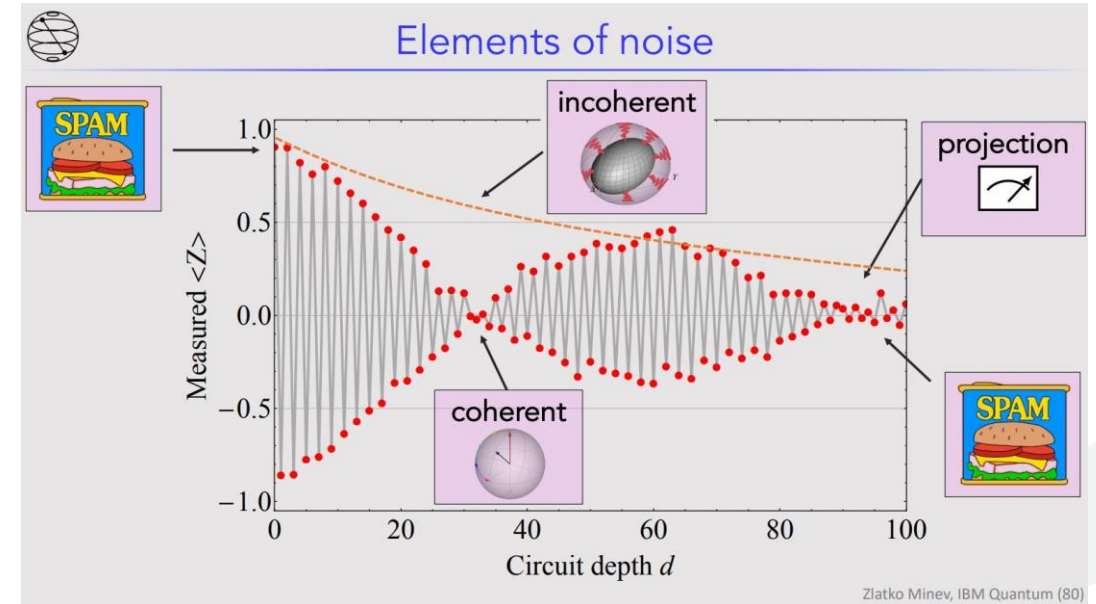
Experimental Results

- Results on Multiple IBM Quantum Computers

CompVQC can reduce circuit length by 2x while the accuracy is also higher in a noisy environment.

Datasets		Syn-Dataset-4		Syn-Dataset-16	
Compression Method		Acc. (vs. Baseline)	TCD (Speedup)	Acc. (vs. Baseline)	TCD (Speedup)
Qiskit Aer	Vanilla VQC	94%(0)	23(0)	96%(0)	51(0)
	Comp-VQC	99%(5%)	11(2.09×)	98%(2%)	23(2.22×)
IBM Q	Vanilla VQC	79%(-15%)	23(1.00×)	86%(-10%)	51(1.00×)
	Comp-VQC	99%(5%)	11(2.09×)	98%(2%)	23(2.22×)

Acc.(vs. Baseline)	ibm_lagos	ibm_perth	ibm_jakarta
Vanilla VQC(TCD=23)	79%(0)	86%(0)	92%(0)
CompVQC(TCD=11)	99%(20%)	98%(12%)	100%(8%)



Circuit compression can make the QNN model more robust to the noise



zwang48@gmu.edu



George Mason University

4400 University Drive
Fairfax, Virginia 22030

Tel: (703)993-1000